

Data Science with **DASK**

2nd MarDATA Block Course
“Advanced Scientific Programming”

November 19th, 2020

Martin Claus, Katharina Höflich, Willi Rath

Agenda for today

- Overview of Dask's fundamental concepts
- First hands-on session on Dask basics
- Review of the key takeaways
- Excursion: Dask distributed deployments / tools
- Second hands-on session on Dask for machine learning
- Question / Answer Session

Let's also schedule breaks?

Aim of this course

We want you to understand the following:

- Dask is all about *task-graphs*.
- Dask provides various ways of *building task graphs* some of which can *replace existing toolboxes* like Numpy or Pandas.
- Dask provides various ways of executing task-graphs in *parallel* on a single machine or on *distributed clusters*.

And: Know when NOT to use Dask!

Dask base concepts

~~Collections~~
(create task graphs)

Task Graph

~~Schedulers~~
(execute task graphs)

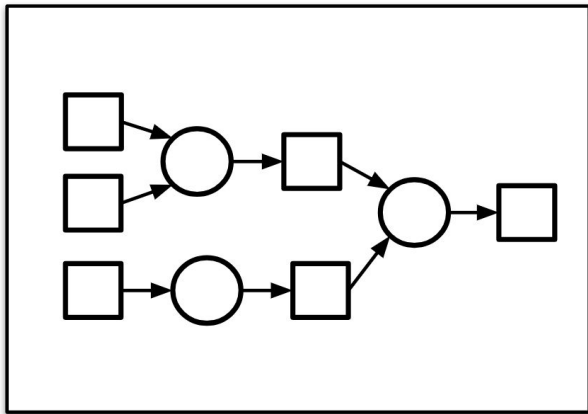
Dask Array

Dask DataFrame

Dask Bag

Dask Delayed

Futures

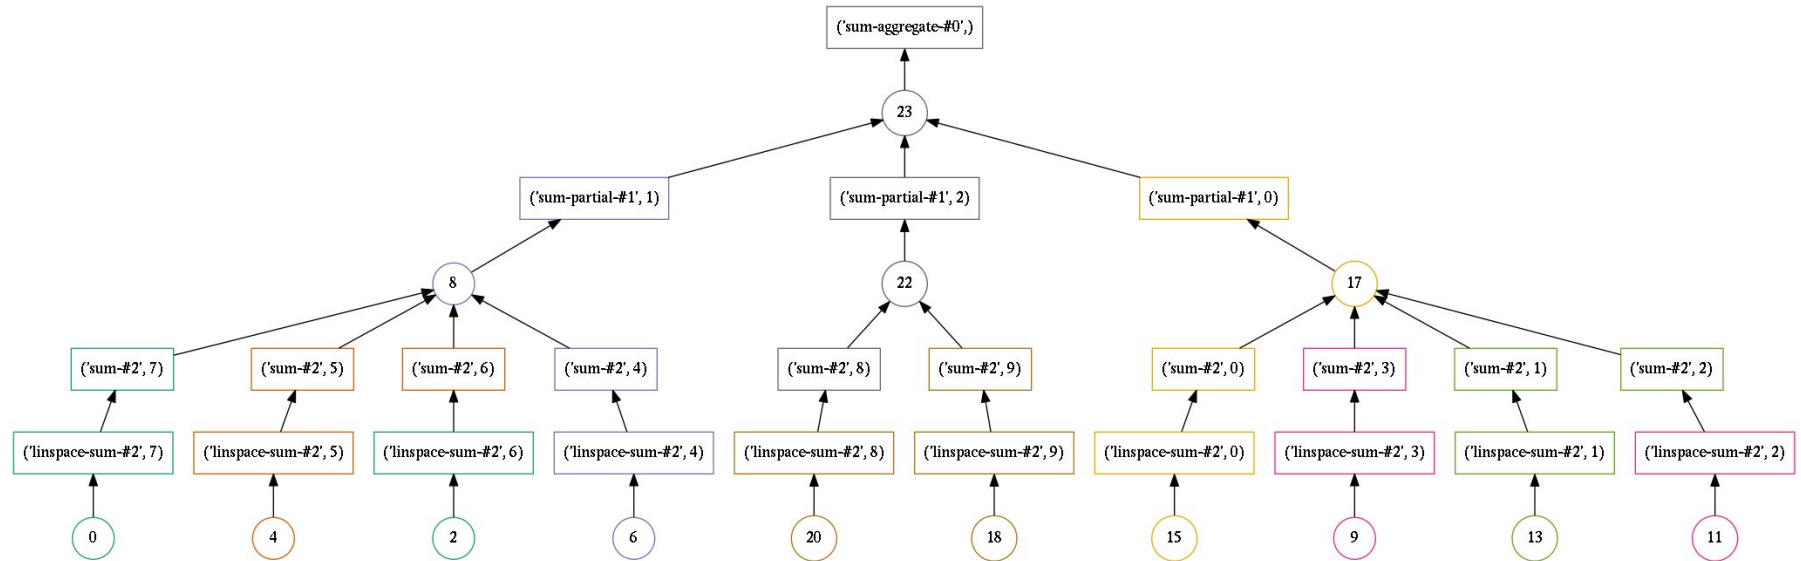


Single-machine
(threads, processes,
synchronous)

Distributed

from: <https://docs.dask.org/en/latest/>

Task graphs encode the flow of information



Your turn: *Can you find workflows / calculations you already know that can be described as a directed acyclic graph?*

Hands on

What is Dask *[30 minutes?]*

- Lecture-style overview of the core concept of Dask.
- Simple examples for explicitly building a task graph.

Creating task graphs *[60 minutes?]*

- Using Dask bag.
- Dask Dataframes.
- Dask Arrays.

Different ways of executing task graphs *[45 minutes?]*

- Get to know a simple application.
- Run the workflow on a single machine.
- Scale `_out_` to multiple machines.

[Lunch ~ 12:30 - 13:00]

Parallelize machine learning with Dask *[45 minutes?]*

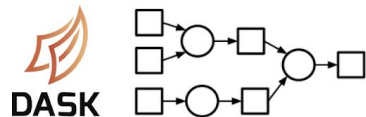
- Parallelize machine learning workloads.

→ <https://github.com/mardatade/Course-Data-Science-with-Dask>

Dask distributed scheduling

- For now, we have only worked with Dask single-machine schedulers.
- The necessary entities to go remote and to scale out are provided by the Dask distributed package (<https://distributed.dask.org/en/latest/>).
- Key strength of Dask: it is designed for the utilization of compute/memory resources of hundreds of separate machines in a common network.
- Let's have a look now... and put together some machines manually!

dask.distributed.Client()



DASK

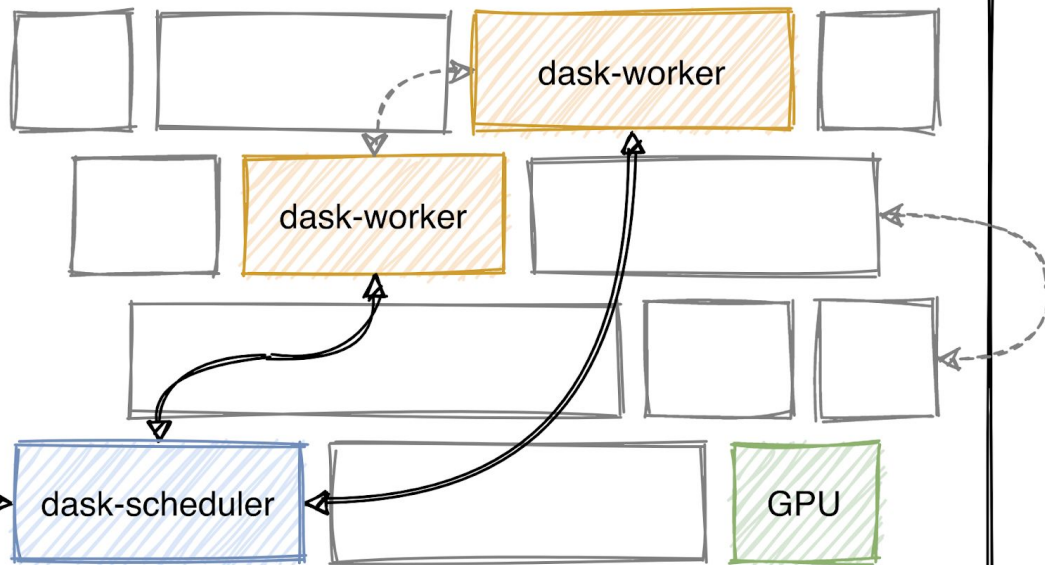
e.g. `dask.array()`



client

connects a remote cluster
to the local namespace that
knows about the task graph

dask.distributed cluster



the scheduler manages
task execution

workers execute tasks, communicate
with each other and store/serve results

Semi-Automatic Ways of Creating Dask Clusters

- For high-performance computers / other compute servers
 - [Dask jobqueue](#): Manages clusters with a single-node granularity. Probably the most mature and most convenient of the HPC deployment solutions.
 - [Dask MPI](#): Uses MPI to distribute scheduler and workers within an HPC job.
(but does not communicate via MPI!)
 - [Dask SSH](#): Starts scheduler and workers on hosts you can access via SSH
 - [Dask DRMAA](#): Uses [DRMAA](#) (a common interface to many different HPC scheduling softwares) for deploying Dask clusters.
- For cloud-computing platforms
 - [Dask kubernetes](#): Creates and manages Dask clusters with Kubernetes
 - [Dask cloudprovider](#): High-level tool that deploys clusters to (Amazon Web Services, Digital Ocean, Google Cloud Platform, Microsoft Azure).
 - [Dask yarn](#): Deploys Dask clusters on Hadoop clusters.
 - [Dask gateway](#): Provides Dask clusters via a web service / API.

... and counting

Summary: Key points to remember

We want you to understand the following:

- Dask is all about *task-graphs*.
- Dask provides various ways of *building task graphs* some of which can *replace existing toolboxes* like Numpy or Pandas.
- Dask provides various ways of executing task-graphs in *parallel* on a single machine or on *distributed clusters*.

And: Know when NOT to use Dask!

What's missing here?

- Dask as backend for other libraries
 - scikit-learn (partially covered here)
 - Xarray (labeled ndimensional arrays): <http://xarray.pydata.org>
 - Workflow management: <https://prefect.io>
- [Actors](#)
- [Optimizing / tuning graphs](#)
- [Debugging](#)
- Input / Output
- Resilience
- ...

Practical Problems

- Task graph should reflect flow of information between functions / methods rather than operations within calculations. Example high-order polynomial with a few coefficients is best wrapped into generalized ufuncs.
- Keep task graphs below a few 100.000 tasks, because scheduling takes time and resources.
- How to choose chunk sizes for Dask arrays?
- ...

Thank you!

What else do you want to know?